




Chapter 31: Information Retrieval


Database System Concepts, 7th Ed.

©Silberschatz, Korth and Sudarshan

See www.db-book.com for conditions on re-use

Introduction to Information Retrieval





Outline

- What is Information Retrieval ?
- Basic IR concepts and models
 - Boolean model
 - Vector space model
 - Probabilistic model
 - Language model

Some of the slides are adapted from the publically available ppt for the textbook “Introduction to Information Retrieval”



Outline

- **What is Information Retrieval ?**
- Basic IR concepts and models

karen sparck jones — karen sparck jones — Google 搜索

karen sparck jones — Google 搜索

你 搜索 图片 地图 Play YouTube 新闻 Gmail 更多

Google karen sparck jones

网页 图片 地图 更多 搜索工具

找到约 34,500 条结果 (用时 0.23 秒)

小提示: 仅取搜索简体中文结果。您可以在设置中指定搜索语言

karen sparck jones 的图片搜索结果 - 举报图片



[Karen Sparck Jones - Wikipedia, the free encyclopedia](#)
en.wikipedia.org/wiki/Karen_Sparck_Jones 翻译此页

www.cl.cam.ac.uk/archive/ksj21. Karen Sparck Jones FBA (26 August 1935 – 4 April 2007) was a British computer scientist. Contents. 1 Personal life; 2 Career.

Education - Career - Further obituaries - References

[Karen Sparck Jones](#)
www.cl.cam.ac.uk/archive/ksj21/ 翻译此页

Karen Sparck Jones passed away in 2007. (Obituary). We pro from her personal web space, in particular those reachable th

[Karen Sparck Jones: 信息检索界最高奖获得者-CSDN](#)
www.csdn.net 业界

2011年3月7日 - Karen Sparck Jones 剑桥大学计算机女科学家... 的言论: 计算机是如此重要, 因此不能把它只留给男人去做!

[Karen Sparck Jones lecture | BCS Academy of Co](#)
academy.bcs.org Our activities - Lectures and events

cn.bing.com/search?q=karen+sparck+jones&go=&qs=n&form=QSLH&pq=karen+sparck+jones&sc=0-17&sp=-1&sk=- karen sparck jones - 必应 Bing

bing - Google 搜索 karen sparck jones - 必应 Bing 搜索搜索引擎 - 上页从使用开始

WEB IMAGES VIDEOS MAPS NEWS DICTIONARY SCORE MORE

必应 karen sparck jones

Tips: Showing results for All English Chinese

[Karen Sparck Jones - Wikipedia, the free encyclopedia](#)
Karen Sparck Jones FBA (26 August 1935 – 4 April 2007) was a British computer scientist. Karen Ida Beal Sparck Jones was born in ... en.wikipedia.org/wiki/Karen_Sparck_Jones 2013-10-18

[Karen Sparck Jones-计算机先驱人物-科技中](#)
... Translate this page
Karen Sparck Jones (1935-2007), 英国剑桥大学计算机女科学家, 英国科学院院士, 微软剑桥研究院前任院长的妻子, Karen Sparck Jones从二十世纪五十年代开 ... www.techcn.com.cn/index.php?doc-view-131747.html 1 day ago

[Images of karen sparck jones](#)
bing.com/images



[dblp: Karen Sparck Jones](#)
1990 - 1999 see FAQ What is the meaning of the colors in the publication list? 1999 [j20] view electronic edition via DOI export record as BibTeX XML dblp key: ... www.informatik.uni-trier.de/~pers/hd/j/Jones_Karen_Sparck 2013-10-19

[Karen Sparck Jones Publications - The Computer Laboratory](#)
Publications M. Masterman, R.M. Needham, K. Sparck Jones and B. Mayoh, "Agricola Incurvo terram dimovit aratro", Cambridge Language Research Unit, Report (ML84 ... www.cl.cam.ac.uk/archive/ksj21/ksjdlgpapers/ksjbb3.html 2010-7-8

[Karen Sparck Jones: 信息检索界最高奖获得者](#)
... Translate this page
摘要: Karen Sparck Jones 剑桥大学计算机女科学家 Karen Sparck Jones 的言论: 计

Ad
[jones双11优惠券](#)
www.tmall.com
积分换取优惠券, 百亿让利提前抢, 锁定心仪品牌, 5折后再抵用, 赶...
See your ad here >

RELATED SEARCHES
[karen jones melbourne florida](#)
[karen](#)
[karen black](#)
[karen mck](#)
[karen chung](#)
[karen walker](#)
[karen是什么歌星](#)
[karen 露口双枪](#)

SEARCH TOOLS
[Turn off Hover Translation](#)



www.baidu.com/s?wd=karen+spark+jones&rsrv_bp=0&ch=&tn=baidu&bar=8rsrv_spt=3&ie=utf-8&rsrv_n=2&rsrv_sug3=1&rsrv_sug4=1&inputT=1056 -- kare...

karen spark jones - Google 搜索 karen spark jones - 必应 Bing karen spark jones_百度搜索

Baidu 百度 新闻 网页 贴吧 知道 音乐 图片 视频 地图 文库 更多

搜索框: karen spark jones 百度一下 推荐: 用手机浏览就上百度

Karen Spark Jones - Wikipedia, the free encyclopedia
 Karen Spark Jones FBA (26 August 1935 – 4 April 2007) was a British computer scientist.[2]
 [3]Contents 1 Education 2 Career 3 Further ...
en.wikipedia.org/wiki/Karen_Spark_Jones 2007-04-04 - 百度快照

DBLP: Karen Spark Jones
 68 EE Karen Spark Jones: Statistics and Retrieval: Past and Future. ICCTA 2007: 396-405
 67 EE Karen Spark Jones: 2007 Athena Lecturer Award ...
ftp.informatik.wth-aachen.de/dblp/... 1995-08-28 - 百度快照

Search results for 'Karen Spark Jones' - PhilPapers
 Search results for 'Karen Spark Jones' (try it on Scholar) 1000+ foundSort by: Karen Jones
 (2006). Metaethics and Emotions Research: A Response to ...
philpapers.org/a/Karen_Spark_Jones 2012-12-18 - 百度快照

Karen Spark Jones: 信息检索界最高奖获得者-CSDN.NET
 2011年3月7日 - Karen Spark Jones 剑桥大学计算机女科学家Karen Spark Jones
 论:计算机是如此重要,因此不能把它只留给男人去做!在程序界广为流传
www.csdn.net/article/2011-03-07/293... 2011-03-07 - 百度快照

Selected Papers of Karen Spark Jones: 亚马逊图书
 "Today, anyone who uses Google to hunt for information on the World Wide Web is aware of the importance of fundamental research conducted by Karen Spark Jones ...
www.amazon.cn/gp/sr/d/1575865... 2004-11-01 - 百度快照

...Systems - Julia Rose Galliers, Karen Spark Jones - Google Books
 作者 Julia Rose Galliers, Karen Spark Jones 投稿人 University of Cambridge Computer Laboratory 出版商 University of Cambridge, Computer Laboratory, 1997
www.google.co.jp/books?id=ZxkIAQAIA... 2010-01-19 - 百度快照

...Karen Spark Jones, S. Walker, S. E. Robertson - Google Books
 Karen Spark Jones, S. Walker, S. E. Robertson 0 评价 University of Cambridge Computer Laboratory, 1998 - 74 页 Abstract: "The paper combines a ...
www.google.com.tw/books?id=fKkdAQAIA... 2009-10-15 - 百度快照



www.sogou.com/web?query=karen+spark+jones&asf=www.sogou.com&ast=1383092497&w=01019900&p=40040100&ie=utf8&sut=2566744&st0=1383092497037&...

karen spark jones - Google 搜索 karen spark jones - 必应 Bing karen spark jones - 搜狗搜索

Sogou 搜狗 新闻 网页 音乐 图片 视频 地图 知识 购物 更多

搜索框: karen spark jones 搜狗搜索

网页 文档 博客 论坛 知识 新闻 更多 什么是分类搜索 找到约 688 条结果

Karen Spark Jones [下载]
 TOWARDS BETTER NLP SYSTEM EVALUATION Karen Spark Jones Computer Laboratory,
 University of Cambridge New Museums Site, Pembroke Street, Cambridge CB2 3QG ...
sci.fdc.upenn.edu/H/H94/H94-1018.pdf - 2011-10-24 - 快照 - 预览

Selected Papers of Karen Spark Jones (豆瓣)
 图书Selected Papers of Karen Spark Jones 介绍、书评、论坛及推荐
 豆瓣读书 - book.douban.com/subject/4674477/ - 2013-7-26 - 快照 - 预览


Sparck Jones, Karen, 1935-2007 - LC Linked Data Service (Library of Congress)
 (hdgs.: Sparck Jones, Karen, 1935-; Jones, K. Sparck (Karen Spark); usage: Karen Spark Jones, K. Sparck Jones, K. Sparck Jones, K.I.B. Sparck Jones)
id.loc.gov/authorities/names/n50021679 - 2005-3-17 - 快照 - 预览

Re Sparck Jones Re Ballard on the Need for Careful Description... (0 财富值)
 大小: 125.18KB 页数: 2
 Re Sparck Jones Re Ballard on the Need for Careful Description s Re Ballard on the Need for Careful Description! read with interest the recent response by Karen Spark Jones ...
wenku.baidu.com/view/445d85... 2011-9-14 - 快照 - 预览

Jones Karen Spark - Download eBooks Audiobooks
 Financial System—and Themselves Andrew Ross Sorkin Rita Mulcahy George R. R. Martin George R. R. Martin George R. R. Martin Roy Dotrice Jones Karen Spark Jones Karen ...
bookstore.com/fr/Jones-Karen-Spark/m68174/ - 2013-10-6 - 快照 - 预览

karen spark jones award-KSJ Award
 Retrieval and Natural Language Processing with significant experimental contributions. The IR SG is grateful to Microsoft Research for their sponsorship of the Karen Spark Jones ...
irsg.bcs.org/ksjaward.php - 2013-4-7 - 快照 - 预览

Karen Spark Jones: 信息检索界最高奖获得者-CSDN.NET



Click to open expanded view



[Share your own customer images](#)

[Search inside this book](#)

Introduction to Information Retrieval [Hardcover]

[Christopher D. Manning](#) (Author), [Prabhakar Raghavan](#) (Author), [Hinrich Schütze](#) (Author)

★★★★☆ (18 customer reviews)

Buy New

\$57.73 & FREE Shipping. [Details](#)

In Stock.

Ships from and sold by Amazon.com. Gift-wrap available.

Want it Thursday, Oct. 31? Order within **19 hrs 40 mins** and choose **One-Day Shipping** at checkout. [Details](#)

52 new from \$29.67 **42 used** from \$28.32

Share

FREE TWO-DAY SHIPPING FOR COLLEGE STUDENTS

[Learn more](#)

amazonstudent

Formats	Amazon Price	New from	Used from
Hardcover	\$57.73	\$28.67	\$28.32

Rent

\$26.00

In Stock.

Rented by [apex media](#) and [Fulfilled by Amazon](#).

Buy New \$57.73

Quantity:

☐ Yes, I want **FREE Two-Day Shipping** with Amazon Prime

Add to Cart

or

[Sign in to turn on 1-Click ordering](#)

☐ **Buy Used \$54.55**

[Add to Wish List](#)

Sell Us Your Item

For a **\$18.72** Gift Card

Trade in

[Learn more](#)

More Buying Choices

94 used & new from \$28.32

Have one to sell? [Sell on Amazon](#)

Save up to **70%**

Rent Your Textbooks

Save up to 70% when you [rent your textbooks](#) on Amazon. Keep your textbook rentals for a semester and rental return shipping is free. [Learn more.](#)

Authors (5599)

Sergios Theodoridis

Hui Jiang

Qiang Huang

Hang Li (李航)

Dawei Song

Conferences (406)

SIGIR

CIKM

ICASSP

WWW

SIGMOD

Journals (203)

TASLP

BMC Bioinformatics

CORR

PR

PAMI

Organizations (1416)

Microsoft

Chinese Academy of Sciences

University of Southern

Academic > Results for "hidden markov model query language model year>=2006" in Computer Science

Subscribe

Publications (2887)

since 2006

[A latent variable model for query expansion using the hidden markov model](#) [View...](#)

Qiang Huang, Dawei Song

...structure of a latent variable model (lvm) for query language modeling. in the proposed lvm, the combinations of query terms are viewed as the...are governed by an ergodic hidden markov model (hmm), where the viterbi algorithm...
Conference: International Conference on Information and Knowledge Management - CIKM, pp. 1417-1418, 2008

[Classification of musical patterns using variable duration hidden Markov models](#) (Citations: 10) [View...](#)

Aggelos Pikrakis, Sergios Theodoridis, Dimitris Karamotou

...a set of variable duration hidden markov models. each of these models has been trained to recognize...new type of variable duration hidden markov model provides increased clas- sification accuracy...
Journal: IEEE Transactions on Audio, Speech & Language Processing - TASLP, vol. 14, no. 5, pp. 1795-1807, 2006

[A Qualitative Hidden Markov Model for Spatio-temporal Reasoning](#) (Citations: 1)

Zina M. Ibrahim, Ahmed Y. Tawfik, Alioune Ngom

...we present a hidden markov model that uses qualitative order of magnitude probabilities for its states and transitions. we use the resulting model to construct a formalization of...and recognition techniques used in hidden markov models to perform useful queries with the representation. ...
Conference: Symbolic and Quantitative Approaches to Reasoning and Uncertainty - ECSQARU, pp. 707-718, 2007

[Keyword query cleaning using hidden Markov models](#) (Citations: 2) [View...](#)

Ken Q. Pu

...consider the problem of keyword query cleaning for structured databases from a probabilistic ap- proach. keyword query cleaning consists of rewriting the user query, segmenting the keywords, matching each...efficient and robust solution using hidden markov models (hmm). by modeling user keyword queries using a generative proba- bilistic hmm-based model, we





首页
热门
应用
游戏
大家正在热搜: 1989
张鹏_天大
✉
✉
⚙
🔒


综合
微博
找人
图片
应用

搜索
高级搜索
设置
帮助



yyuegguang: 请饶过农民吧——“秸秆燃烧产生的烟雾通过平流作用, 客观上会对**京津冀地区雾霾**的加重起到助推作用, 但秸秆燃烧并不是**雾霾**产生的主要原因。污染排放和天气因素的共同作用是霾形成和加重的原因, 解决霾污染的根本途径在于减排。” <http://t.cn/R7Jv01S>



3分钟前 来自分享按钮

👍 | 转发 | 收藏 | 评论

★ 精选



人民日报 : 【16座“洗肺”城市 有你的家乡吗?】**京津冀雾霾**再度来袭, 拉仇恨的是, 有些城市却不必遭遇“十面霾伏”。9月最新数据, 这16个城市空气质量位居全国前列: 海口、舟山、拉萨、盐城、哈尔滨、福州、厦门、贵阳、张家口、惠州、台州、昆明、连云港、丽水、江门、深圳。羡慕吗? 计划一次“洗肺”之旅吧?






找人结果 更多>>



京津冀天空无雾霾
简介: 无所畏与无所谓。快乐工作, 享受生活。做个快乐的人。...




京津冀雾霾一体化
简介: 京津冀一体化的速度, 赶不上雾霾的脚步!


实时搜索热点 更多>>

1. [nikestore](#) (15794)
2. [1989](#) (8819)
3. [夏语心](#) (8538)
4. [马蓉](#) (6045)
5. [西蒙子](#) (5874)
6. [林更新](#) (5768)
7. [唐嫣](#) (5336)
8. [郑凯](#) (4983)
9. [庞麦郎](#) (2984)
10. [跑男](#) (2964)

我的订阅词 订阅当前词

你关注的词 (0/10)





Common features of information retrieval applications



- **Goal:** To retrieve information that is **relevant** to the user's **information need** and help the user complete an information seeking **task**


Information need:

books, personal info, products, news



Information Retrieval

- The subject that aims at seeking desired information given a user's **information need**. It involves information acquisition, storage, organization and access.
- The process of finding **material** of an **unstructured nature** (usually text or multimedia) that **satisfies an information need** from within large collections (usually stored on computers).



Information Retrieval

- How to find relevant objects?

Define and calculate the “**similarity**” function:


- Give the information need: **query**
- Objects to be retrieved: **text or multimedia**
- Match between query and object: **relevance score**



The classical search model

Information Object

Information Need





IR is also science

- IR is not merely an engineering problem, but also a scientific subject
 - NLP
 - Pattern recognition,
 - Artificial intelligence
 - Machine learning
 - Human-machine interaction
 - User behavior and cognition
 - Modeling complex information objects
 - Social networks
 - Big data
 -

The complex interactions among above topics



IR Applications

- Web search engines (e.g., Google, Baidu, Bing, Sougou)
- Intranet and domain-specific search systems
- E-commerce (e.g., Alibaba, Amazon)
- Social networking site (e.g., Weibo, Facebook, Twitter)
- Intelligence and security
- Mobile apps
- Digital libraries
- Consumer behavior analysis
- News/Ads push
- Question-Answering system
-



Outline


- What is Information Retrieval ?
- **Basic IR concepts and models**



The classic search model

Information Object

Information Need





Basic IR concepts

- The process of finding **material** of an **unstructured nature** (usually text or multimedia) that **satisfies an information need** within large collections (usually stored on computers).
- **Unstructured**
- **Documents**
- **Information Need**
- **Collection**

IR vs. databases:

Structured vs unstructured data



- Structured data tends to refer to information in “tables”

Employee	Manager	Salary
Smith	Jones	50000
Chang	Smith	60000
Ivy	Smith	50000

Typically allows numerical range and exact match (for text) queries, e.g.,

Salary < 60000 AND Manager = Smith.

IR vs. databases



- IR systems don't deal with transactional updates (including concurrency control and recovery)
- Database systems deal with structured data, with schemas that define the data organization
- IR systems deal with some querying issues not generally addressed by database systems
 - Approximate searching by keywords
 - Ranking of retrieved answers by estimated degree of relevance



Unstructured data

- Typically refers to free text
- Allows
 - Keyword queries including operators
 - More sophisticated queries e.g.,
 - find all web pages dealing with *drug abuse*
- Classic models for searching textual documents



Semi-structured data

- In fact almost no data is “unstructured”
<title>homepage</title>
<body>...</body>
- Facilitates “semi-structured” search such as
 - *Title* contains data AND *Bullets* contain search
 - other structures not mentioned here, such as linguistics and Hyperlinks



Assumptions of classical IR

- **Collection:** A set of documents
 - Assume it is a static collection for the moment
- **Goal:** Retrieve documents with information that is **relevant** to the user's **information need** and helps the user complete a **task**
- **Relevance:**
 - Objective
 - The degree of matching
 - Depend on specific applications

Term-document incidence matrices




	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	1	1	0	0	0	1
Brutus	1	1	0	1	0	0
Caesar	1	1	0	1	1	1
Calpurnia	0	1	0	0	0	0
Cleopatra	1	0	0	0	0	0
mercy	1	0	1	1	1	1
worser	1	0	1	1	1	0

Brutus AND Caesar BUT NOT Calpurnia

1 if play contains
word, 0 otherwise

Bigger collections



- Consider $N = 1$ million documents, each with about 1000 words.
- Say there are $M = 500K$ *distinct* terms among these.




Can't build the matrix


- 500K x 1M matrix has half-a-trillion 0's and 1's.
- But it has no more than one billion 1's.
 - matrix is extremely sparse.
- What's a better representation?
 - We only record the 1 positions.



Inverted index

- For each term t , we must store a list of all documents that contain t .
 - Identify each doc by a **docID**, a document serial number





Indexer steps: Token sequence

- Sequence of (Modified token, Document ID) pairs.

Doc 1

I did enact Julius
Caesar I was killed
i' the Capitol;
Brutus killed me.

Doc 2

So let it be with
Caesar. The noble
Brutus hath told you
Caesar was ambitious



Term	docID
I	1
did	1
enact	1
julius	1
caesar	1
I	1
was	1
killed	1
i'	1
the	1
capitol	1
brutus	1
killed	1
me	1
so	2
let	2
it	2
be	2
with	2
caesar	2
the	2
noble	2
brutus	2
hath	2
told	2
you	2
caesar	2
was	2
ambitious	2



Indexer steps: Dictionary & Postings


- Multiple term entries in a single document are merged.
- Split into Dictionary and Postings
- Doc. frequency information is added.

Core indexing step

Term	docID
ambitious	2
be	2
brutus	1
brutus	2
capitol	1
caesar	1
caesar	2
caesar	2
did	1
enact	1
hath	1
i	1
i	1
i'	1
it	2
julius	1
killed	1
killed	1
let	2
me	1
noble	2
so	2
the	1
the	2
told	2
you	2
was	1
was	2
with	2



term	doc. freq.	→	postings lists
ambitious	1	→	2
be	1	→	2
brutus	2	→	1 → 2
capitol	1	→	1
caesar	2	→	1 → 2
did	1	→	1
enact	1	→	1
hath	1	→	2
i	1	→	1
i'	1	→	1
it	1	→	2
julius	1	→	1
killed	1	→	1
let	1	→	2
me	1	→	1
noble	1	→	2
so	1	→	2
the	2	→	1 → 2
told	1	→	2
you	1	→	2
was	2	→	1 → 2
with	1	→	2




IR models

- Assume the inverted index has been build up.
- How do we process a query to match with documents to produce the relevance score?
 - Boolean model
 - Vector space model
 - Probabilistic model
 - Language model
 -



Boolean retrieval model

- Consider processing the query (using AND):
Brutus AND Caesar
 - Locate ***Brutus*** in the Dictionary;
 - Retrieve its postings.
 - Locate ***Caesar*** in the Dictionary;
 - Retrieve its postings.
 - “Merge” the two postings (intersect the document sets):





Other Boolean expressions

- **OR:** Brutus OR Caesar
 - The union of two inverted doc sets
- **NOT:** Brutus AND NOT Caesar
 - The difference of two inverted doc sets
- General Boolean expressions
 - (Brutus OR Caesar) AND NOT (Antony OR Cleopatra)

Advantages of Boolean retrieval model



- Easy to establish, perhaps the simplest model to build an IR system upon
 - Primary commercial retrieval tool for 3 decades
 - Many professional searchers still like Boolean search – e.g., medical search
 - You know exactly what you are getting
 - But that doesn't mean it actually works better



Problem with Boolean search

- Boolean expression can be very complex
 - It needs a lot of skills to come up with a query that produces a manageable number of hits.
 - AND gives too few; OR gives too many
 - Only good for expert users with precise understanding of their needs and the collection.
- It does not use the frequency of terms
- Documents either match or don't. It can not order over the documents in the collection for a query



Ranked retrieval models


- In **ranked retrieval**, the system returns an ordering over the (top) documents in the collection for a query
- Produces a manageable number of hits
- **Free text queries:** Rather than a query language of operators and expressions, the user's query is just one or more words in a human language
- **Precondition:** The ranking algorithm is valid, i.e., more relevant documents are ranked before less relevant ones.

Scoring as the basis of ranked retrieval




- We wish to return in order the documents most likely to be useful to the searcher
 - How can we rank-order the documents in the collection with respect to a query?
- Assign a score: say in $[0, 1]$, to each document
 - This score measures how well document and query “match”.

Query-document matching scores



- We need a way of assigning a score to a query-document pair
 - Let's start with a one-term query
 - If the query term does not occur in the document: score should be 0
 - The more frequent the query term in the document, the higher the score (should be)
- We will look at a number of alternatives for this.



Jaccard coefficient

- A commonly used measure of overlap of two sets A and B
 - $\text{jaccard}(A,B) = |A \cap B| / |A \cup B|$
 - $\text{jaccard}(A,A) = 1$
 - $\text{jaccard}(A,B) = 0$ if $A \cap B = 0$
- A and B don't have to be the same size.
- Always assigns a number between 0 and 1.


Jaccard coefficient: Scoring example



- Query: *ides of march*
- Document 1: *caesar died in march*
- Document 2: *the long march*

- Jaccard($q, d1$) = 1/6

Issues with Jaccard for scoring




- It doesn't consider *term frequency* (how many times a term occurs in a document)
- Rare terms in a collection are more informative than frequent terms. Jaccard doesn't consider this information
- We need a more sophisticated way of normalizing for length

Binary term-document incidence matrix



	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	1	1	0	0	0	1
Brutus	1	1	0	1	0	0
Caesar	1	1	0	1	1	1
Calpurnia	0	1	0	0	0	0
Cleopatra	1	0	0	0	0	0
mercy	1	0	1	1	1	1
worser	1	0	1	1	1	0

Each document is represented by a binary vector $\in \{0,1\}^{|V|}$



Term-document count matrices

- Consider the number of occurrences of a term in a document:
 - Each document is a count vector in $\mathbb{N}^v$: a column below

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	157	73	0	0	0	0
Brutus	4	157	0	1	0	0
Caesar	232	227	0	2	1	1
Calpurnia	0	10	0	0	0	0
Cleopatra	57	0	0	0	0	0
mercy	2	0	3	5	5	1
worser	2	0	1	1	1	0



Term frequency tf

- The term frequency $tf_{t,d}$ of term t in document d is defined as the number of times that t occurs in d .
- Raw term frequency is not what we want:
 - A document with 10 occurrences of the term is more relevant than a document with 1 occurrence of the term.
 - But not 10 times more relevant.
- Relevance does not increase proportionally with term frequency.




Log-frequency weighting

- The log frequency weight of term t in d is

$$w_{t,d} = \begin{cases} 1 + \log_{10} \text{tf}_{t,d}, & \text{if } \text{tf}_{t,d} > 0 \\ 0, & \text{otherwise} \end{cases}$$

- $0 \rightarrow 0, 1 \rightarrow 1, 2 \rightarrow 1.3, 10 \rightarrow 2, 1000 \rightarrow 4$, etc.
- Score for a document-query pair: sum over terms t in both q and d
- The score is 0 if none of the query terms is present in the document.



Weighting for rare terms

- Rare terms are more informative than frequent terms
 - Recall stop words
- Consider a term in the query that is rare in the collection (e.g., *arachnocentric*)
 - A document containing this term is very likely to be relevant to the query “*arachnocentric*”
 - We want a high weight for rare terms like *arachnocentric*.
 - Inverse Document Frequency (IDF)




idf weight

- df_t is the document frequency of t : the number of documents that contain t
 - df_t is an inverse measure of the informativeness of t
 - $df_t \leq N$
- We define the idf (inverse document frequency) of t by

$$idf_t = \log_{10} (N/df_t)$$

- We use $\log (N/df_t)$ instead of N/df_t to “dampen” the effect of idf.
- Idf is an measure for the informativeness of term t

idf example, suppose $N = 1$ million



term	df_t	idf_t
calpurnia	1	6
animal	100	4
sunday	1,000	3
fly	10,000	2
under	100,000	1
the	1,000,000	0

$$idf_t = \log_{10} (N/df_t)$$

There is one idf value for each term t in a collection.

Collection vs. Document frequency




- The collection frequency of t is the number of occurrences of t in the collection, counting multiple occurrences.
- Example:

Word	Collection frequency	Document frequency
<i>insurance</i>	10440	3997
<i>try</i>	10422	8760

- Which word is a better search term (and should get a higher weight)?
- df (and idf) is more suitable as weight, compared with cf (and icf)

tf-idf weighting



- The tf-idf weight of a term is the product of its tf weight and its idf weight.

$$w_{t,d} = \log(1 + \text{tf}_{t,d}) \times \log_{10}(N / \text{df}_t)$$

- **Best known weighting scheme in information retrieval**
 - Note: the “-” in tf-idf is a hyphen, not a minus sign!
 - **Alternative names: tf.idf, tf x idf**
- Increases with the number of occurrences within a document
- **Increases with the rarity of the term in the collection**

Binary $\rightarrow$ count $\rightarrow$ weight matrix



	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	5.25	3.18	0	0	0	0.35
Brutus	1.21	6.1	0	1	0	0
Caesar	8.59	2.54	0	1.51	0.25	0
Calpurnia	0	1.54	0	0	0	0
Cleopatra	2.85	0	0	0	0	0
mercy	1.51	0	1.9	0.12	5.25	0.88
worser	1.37	0	0.11	4.15	0.25	1.95


Each document is now represented by a real-valued vector of tf-idf weights $\in \mathbb{R}^{|V|}$



Documents as vectors

- So we have a $|V|$ -dimensional vector space
- Terms are axes of the space
- Documents are points or vectors in this space
- Very high-dimensional: tens of millions of dimensions when you apply this to a web search engine
- These are very sparse vectors - most entries are zero.

Queries as vectors



- Represent queries as vectors in the space
- Rank documents according to their similarity to the query in this space

cosine(query,document)



$$\cos(\vec{q}, \vec{d}) = \frac{\vec{q} \bullet \vec{d}}{|\vec{q}| |\vec{d}|} = \frac{\vec{q}}{|\vec{q}|} \bullet \frac{\vec{d}}{|\vec{d}|} = \frac{\sum_{i=1}^{|V|} q_i d_i}{\sqrt{\sum_{i=1}^{|V|} q_i^2} \sqrt{\sum_{i=1}^{|V|} d_i^2}}$$

The equation is annotated with two boxes: "Dot product" pointing to the dot product term $\vec{q} \bullet \vec{d}$ in the first fraction, and "Unit vectors" pointing to the normalized vectors $\frac{\vec{q}}{|\vec{q}|}$ and $\frac{\vec{d}}{|\vec{d}|}$ in the second fraction.


q_i is the tf-idf weight of term i in the query

d_i is the tf-idf weight of term i in the document

$\cos(\vec{q}, \vec{d})$ is the cosine similarity of $\vec{q}$ and $\vec{d}$... or, equivalently, the cosine of the angle between $\vec{q}$ and $\vec{d}$.



Cosine similarity illustrated





Cosine similarity amongst 3 documents

How similar are
the novels

SaS: *Sense and
Sensibility*


PaP: *Pride and
Prejudice*, and

WH: *Wuthering
Heights*?

term	SaS	PaP	WH
affection	115	58	20
jealous	10	7	11
gossip	2	0	6
wuthering	0	0	38

Term frequencies (counts)

Note: To simplify this example, we don't do idf weighting.



3 documents example contd.

Log frequency weighting

term	SaS	PaP	WH
affection	3.06	2.76	2.30
jealous	2.00	1.85	2.04
gossip	1.30	0	1.78
wuthering	0	0	2.58

After length normalization

term	SaS	PaP	WH
affection	0.789	0.832	0.524
jealous	0.515	0.555	0.465
gossip	0.335	0	0.405
wuthering	0	0	0.588

$$\cos(\text{SaS}, \text{PaP}) \approx 0.789 \times 0.832 + 0.515 \times 0.555 + 0.335 \times 0.0 + 0.0 \times 0.0 \approx 0.94$$

$$\cos(\text{SaS}, \text{WH}) \approx 0.79$$

$$\cos(\text{PaP}, \text{WH}) \approx 0.69$$


Why do we have $\cos(\text{SaS}, \text{PaP}) > \cos(\text{SaS}, \text{WH})$?

Summary – vector space ranking



- Represent the query as a weighted tf-idf vector
- Represent each document as a weighted tf-idf vector
- Compute the cosine similarity score for the query vector and each document vector
- Rank documents with respect to the query by score
- Return the top K (e.g., $K = 10$) to the user

Relevance Using Hyperlinks



- Using term frequencies makes “spamming” easy
 - E.g., a travel agency can add many occurrences of the words “travel” to its page to make its rank very high
- Most of the time people are looking for pages from popular sites
- Idea: use popularity of Web site (e.g., how many people have visited it) to rank site pages that match given keywords
- Problem: hard to find actual popularity of site
 - Solution: next slide



Relevance Using Hyperlinks (Cont.)

- Solution: use number of hyperlinks to a site as a measure of the popularity or **prestige** of the site
 - Count only one hyperlink from each site
 - Popularity measure is for site, not for individual page
 - But, most hyperlinks are to root of site
 - Also, concept of “site” difficult to define since a URL prefix like `cs.yale.edu` contains many unrelated pages of varying popularity
- Refinements
 - When computing prestige based on links to a site, give more weight to links from sites that themselves have higher prestige
 - Definition is circular
 - Set up and solve system of simultaneous linear equations
 - Above idea is basis of the Google **PageRank** ranking mechanism

Relevance Using Hyperlinks (Cont.)




- Connections to **social networking** theories that ranked prestige of people
 - E.g., the president of a country has a high prestige since many people know him/her
 - Someone known by multiple prestigious people has high prestige
- Hub and authority based ranking
 - A **hub** is a page that stores links to many pages (on a topic)
 - An **authority** is a page that contains actual information on a topic
 - Each page gets a **hub prestige** based on prestige of authorities that it points to
 - Each page gets an **authority prestige** based on prestige of hubs that point to it
 - Again, prestige definitions are cyclic, and can be got by solving linear equations
 - Use authority prestige when ranking answers to a query



- **Probabilistic model**
- **Statistical Language model**



Why probabilities in IR?



In traditional IR systems, matching between each document and query is attempted in a semantically imprecise space of index terms.

Probabilities provide a principled foundation for uncertain reasoning.
Can we use probabilities to quantify our uncertainties?



Probabilistic IR models

- **Ranking method is the core of an IR system:**
 - In what order do we present documents to the user?
- **Prob-IR is to rank by probability of relevance of the document w.r.t. information need**
 - $R=\{0,1\}$ denotes the relevance judgment,
 - $Q=\{q_1, q_2, \dots\}$ is the query,
 - $D=\{d_1, d_2, \dots\}$ represents the document.
 - Then the conditional probability $P(R=1|Q=q, D=d)$ denotes the relevance score w.r.t. the document and query.

Examples of Probabilistic IR model



- **Binary independence model (BIM)**
- **BM25 (BestMatch25)**
- **Language Model**

Binary Independence Model (BIM)



- **BIM**: Robertson and Sparck Jones, Presented in 1970s, the OKAPI system
- Bayes' Rule

$$P(A | B) = \frac{P(A, B)}{P(B)} = \frac{P(B | A)P(A)}{P(B)}$$

- Based on Bayes' rule, BIM estimates the conditional probability $P(R=1|Q,D)$. BIM is a generative model.
- For the same Q , $P(R=1|Q,D)$ is denoted as $P(R=1|D)$

$$P(R=1|D) = P(D|R=1)P(R=1) / P(D)$$



BIM (continued)

- For each Q , *define the* Ranking function: Retrieval Status Value $RSV(Q,D)$:

$$\log \frac{P(R=1|D)}{P(R=0|D)} = \log \frac{P(D|R=1) \cancel{P(R=1)/P(D)}}{P(D|R=0) \cancel{P(R=0)/P(D)}}$$

$\propto \log \frac{P(D|R=1)}{P(D|R=0)}$

A constant for each query

where, $P(D|R=1)$, $P(D|R=0)$ represent the generative probabilities of D on the condition of relevance/non-relevance respectively.

$P(D|R=1)$: If a relevant document is retrieved, how probably it will be D ?



How to generate a document?

- Probabilistic viewpoint:
 - Terms follow a certain distribution. We draw samples (terms) from that distribution, which are assembled together to form a document.
 - For $P(D|R=1)$ or $P(D|R=0)$, we assume that $R=1$ or 0 denotes different term distributions, from which the document is generated.

Two common distribution families to generate documents



- **Multi-variate Bernoulli distribution**

- M is the size of the word dictionary, and each term corresponds to one biased coin (head or tail). The probability of being head for the i -th coin is p_i
- Assume: $M=4$ (I you can fly), $p_1=0.7$, $p_2=0.4$, $p_3=0.1$, $p_4=0.05$
- Then, $P(\text{I can fly fly})=0.7*(1-0.4)*0.1*0.05$

- **Multinomial distribution**

- A biased dice with M faces, and each term corresponds to one face. The probability of the upper face to be the i -th face is p_i
- Assume: $M=4$ (I you can fly), $p_1=0.4$, $p_2=0.3$, $p_3=0.2$, $p_4=0.1$
- Then, $P(\text{I can fly fly})=0.4*0.2*0.1*0.1$

Estimate the $P(D|R=1)$ and $P(D|R=0)$ in BIM



- **Analogy: M independence experiments (Multi-variate Bernoulli distribution)**
 - Assume M terms, i.e., M biased coins. The i -th coin: head represents the occurrence of t_i and tail represents the non-occurrence of t_i . Toss the M coins, and records the coins with head face-up. Then the corresponding terms are combined to form document D .
 - Therefore, $P(D|R)$ is indeed the probability of document D by tossing M coins. Assume different coin tosses are independent of each other (**the independence assumption**). Also assume that we only consider the occurrence/non-occurrence of t_i (**binary**), instead of the frequency of t_i . The order of coin tosses is neglected (**bag-of-words**).



An example

- Query: 信息(information) 检索(Retrieval) 教程 (course)
All the term probabilities conditioned on relevance and irrelevance p_i 、 q_i are:

词项	信息	检索	教材	教程	课件
R=1时的概率 p_i	0.8	0.9	0.3	0.32	0.15
R=0时的概率 q_i	0.3	0.1	0.35	0.33	0.10

Document D1: 检索 (retrieval) 课件 (ppt)

$$P(D|R=1)=(1-0.8)*0.9*(1-0.3)*(1-0.32)*0.15$$

$$P(D|R=0)= (1-0.3)*0.1*(1-0.35)*(1-0.33)*0.10$$

$$P(D|R=1)/P(D|R=0)=4.216$$



The derivation of BIM

Let the document D be $\bigwedge_{t_i \in D} t_i \bigwedge_{t_j \notin D} \bar{t}_j$

$$P(D | R = 0) = \prod_{t_i \in D} P(t_i | R = 0) \prod_{t_i \notin D} P(\bar{t}_i | R = 0)$$

$$= \prod_{t_i} q_i^{e_i} (1 - q_i)^{1-e_i}, \text{ if } t_i \in D \text{ then } e_i = 1, \text{ else } e_i = 0$$

$$P(D | R = 1) = \prod_{t_i \in D} P(t_i | R = 1) \prod_{t_i \notin D} P(\bar{t}_i | R = 1)$$

$$= \prod_{t_i} p_i^{e_i} (1 - p_i)^{1-e_i}, \text{ if } t_i \in D \text{ then } e_i = 1, \text{ else } e_i = 0$$

$$\begin{aligned} p_i &= P(t_i | R = 1) \\ q_i &= P(t_i | R = 0) \end{aligned}$$

Note: $P(t_i|R=1)$ denotes the probability of the occurrence of t_i in a document conditioned on the relevance assumption (for some cases, $P(t_i|R=1)$ could be 1). This is not the probability of t_i that occurs in the whole relevant documents. Thus $P(t_i|R=1)$ does not sum up to 1.



The derivation of BIM

- Remove the constant component that only depends on Q. Also assume that for non-match terms in Q, we have $p_i=q_i$.

$$\log \frac{P(D | R=1)}{P(D | R=0)} = \log \frac{\prod_{t_i \in D \cup \bar{D}} p_i^{e_i} (1-p_i)^{1-e_i}}{\prod_{t_i \in D \cup \bar{D}} q_i^{e_i} (1-q_i)^{1-e_i}} = \sum_{t_i \in D \cup \bar{D}} \log \left(\frac{p_i}{q_i} \right)^{e_i} \left(\frac{1-p_i}{1-q_i} \right)^{1-e_i} \quad \text{constant}$$

The weight of t_i in D is 0 or 1

$$= \sum_{t_i \in D \cup \bar{D}} \left(e_i \log \frac{p_i}{q_i} + (1-e_i) \log \frac{1-p_i}{1-q_i} \right) = \sum_{t_i \in D \cup \bar{D}} \left(e_i \log \frac{p_i}{q_i} - e_i \log \frac{1-p_i}{1-q_i} + \log \frac{1-p_i}{1-q_i} \right)$$

Similar to the inner product in vector space

$$\propto \sum_{t_i \in D \cup \bar{D}} e_i \log \frac{p_i / (1-p_i)}{q_i / (1-q_i)} = \sum_{t_i \in D} \log \frac{p_i / (1-p_i)}{q_i / (1-q_i)} = \sum_{t_i \in Q \cap D} \log \frac{p_i / (1-p_i)}{q_i / (1-q_i)} + \sum_{t_i \notin Q \wedge t_i \in D} \log \frac{p_i / (1-p_i)}{q_i / (1-q_i)}$$

The Weight of t_i in Q, only depends on Q.

Non-match terms in Q, $p_i=q_i$, This component is zero.

$$\approx \sum_{t_i \in Q \cap D} \log \frac{p_i / (1-p_i)}{q_i / (1-q_i)} = \sum_{t_i \in D \cap Q} W_i^{BIM}$$

The original BIM model, the key is to calculate p_i, q_i



Calculation of p_i and q_i

Ideally, we could separate the whole collection (the document number is known) into four parts based on the relevance to query and the occurrence of t_i :


	relevant R_i (100)	not relevant $N-R_i$ (400)
t_i present n_i (200)	r_i (35)	$n_i - r_i$ (165)
t_i absent $N-n_i$ (300)	$R_i - r_i$ (65)	$N - R_i - n_i + r_i$ (235)

Where N , n_i represent the total number of documents and those documents that t_i occurs in. R_i , r_i represent the number of relevant documents and those that contain t_i . Then:

$$p_i = \frac{r_i}{R_i} = \frac{35}{100} = 0.35$$

$$q_i = \frac{n_i - r_i}{N - R_i} = \frac{165}{400} = 0.413$$

With smoothing



$$p_i = \frac{r_i + 0.5}{R_i + 0.5}$$

$$q_i = \frac{n_i - r_i + 0.5}{N - R_i + 0.5}$$




RSJ weight

- Robertson & Spärck Jones weight (RSJ weight)

$$W_i^{RSJ} = \log \frac{(r_i + 0.5)(N - R - n_i + r_i + 0.5)}{(n_i - r_i + 0.5)(R - r_i + 0.5)}$$

Okapi BM25



- BIM is the simplest way to rank documents:

$$RSV(Q, D) = \sum_{t_i \in D \cup Q} W_i^{IDF}$$

- Consider term frequency, we have:

$$RSV(Q, D) = \sum_{t_i \in D \cup Q} W_i^{IDF} \frac{(k_1 + 1)tf_{t_i, D}}{k_1((1 - b) + b \times (L_D / L_{ave})) + tf_{t_i, D}}$$

- **tf_{ti,D}**: frequency of ti in document D
- **L_D (L_{ave})**: length of document D (averaged document length)
- **k₁**: parameter to control the weight of term frequency
- **b**: parameter to control the weight of document length

Okapi BM25



- For long query, introduce the tf in Q:

$$RSV(Q, D) = \sum_{t_i \in D \cup Q} W_i^{IDF} \cdot \frac{(k_1 + 1)tf_{t_i, D}}{k_1((1 - b) + b \times (L_D / L_{ave})) + tf_{t_i, D}} \cdot \frac{(k_3 + 1)tf_{t_i, Q}}{k_3 + tf_{t_i, Q}}$$

- **tf_{ti, Q}**: the term frequency of ti in Q
- **k₃**: parameter to control the weight of term frequency in query.
- No normalization w.r.t. query length
- Ideally, the above parameters should be tuned to the best performance. Experiments show that k_1 and k_3 should be some value between 1.2 and 2, and b be 0.75.



Statistical Language Model (SLM)

Statistical Language Modeling, SLM



- SLM is widely applied in speech recognition and statistical machine translation. The properties of natural language are analyzed using probability theory.
 - Rule-based: word, sentence, paragraph are generated by some rules.
 - Statistical method: any language pieces are possible. The difference is the probability of occurrence
- For a sequence of terms in the document $d=w_1w_2...w_n$, SLM calculates the probability $P(w_1w_2...w_n)$:

$$P(w_1w_2...w_n) = P(w_1)P(w_2...w_n | w_1) = ... = P(w_1) \prod_{i=2}^n P(w_i | w_{i-1}...w_1)$$

history

No history, Unigram

Nearest one history, Bigram

Nearest N-1 histories, N-gram⁹²



Examples for different SLM

- unigram: $P(w_1 w_2 w_3 w_4) = P(w_1)P(w_2)P(w_3)P(w_4)$
- bigram: $P(w_1 w_2 w_3 w_4) = P(w_1)P(w_2 | w_1)P(w_3 | w_2)P(w_4 | w_3)$
 - 1-order Markov chain
- trigram: $P(w_1 w_2 w_3 w_4) = P(w_1)P(w_2 | w_1)P(w_3 | w_1 w_2)P(w_4 | w_2 w_3)$
- For n-gram, the model complexity grows with greater n. And there are more parameters to be estimated. Of course, when the dataset is sufficiently large, the higher-order model is more accurate.

Similar to strategy in Poker game



Determining which
card to play based
on:

current turn only,
unigram;

previous turn, bi-
gram;

previous 2 turns,
tri-gram

.....



An example of SLM

w	$P(w q_1)$	w	$P(w q_1)$
STOP	0.2	toad	0.01
the	0.2	said	0.03
a	0.1	likes	0.02
frog	0.01	that	0.04
		...	...

unigram

where STOP is a sign for the termination.

The probability for generating the following sentence:

$$P(\text{frog said that toad likes frog STOP}) = 0.01 \cdot 0.03 \cdot 0.04 \cdot 0.01 \cdot 0.02 \cdot 0.01 \cdot 0.02 = 0.00000000000048$$



Two different SLMs

language model of d_1

w	$P(w .)$	w	$P(w .)$
STOP	.2	toad	.01
the	.2	said	.03
a	.1	likes	.02
frog	.01	that	.04
...	...	...	...

language model of d_2

w	$P(w .)$	w	$P(w .)$
STOP	.2	toad	.02
the	.15	said	.03
a	.08	likes	.02
frog	.01	that	.05
...	...	...	...

Q = frog said that toad likes frog STOP

$$P(Q|M_{d_1}) = 0.01 \cdot 0.03 \cdot 0.04 \cdot 0.01 \cdot 0.02 \cdot 0.01 \cdot 0.02 = 0.00000000000048 = 4.8 \cdot 10^{-12}$$

$$P(Q|M_{d_2}) = 0.01 \cdot 0.03 \cdot 0.05 \cdot 0.02 \cdot 0.02 \cdot 0.01 \cdot 0.02 = 0.00000000000120 = 12 \cdot 10^{-12}$$
$$P(Q|M_{d_1}) < P(Q|M_{d_2})$$

Therefore, compare with d_1 , document d_2 is more relevant to “frog said that toad likes frog STOP” .



SLM-based IR model (SLMIR)

- **Query likelihood model:** define the relevance as the generative probability of the current query w.r.t. each document.
- **Translation model:** define the relevance as the probability that the query would have been generated as a translation of the document, and factor in the user's general preferences in the form of a prior distribution over documents.
- **KL-divergence model:** query and document are correspond to two different languages. And the relevance is defined as the KL-divergence between the two language models.
- **We focus on query likelihood model in this lecture.**

Population distribution and sampling



- The document's language model is some population distribution.
- Both document and query are **samples** from that distribution
- Determine the document language model based on the document, i.e., estimate the population distribution.
- Calculate the probability to draw the query from the population distribution



Query Likelihood Model

- Derivation of QLM:

$$P(D|Q) = \frac{P(Q|D)P(D)}{P(Q)} \propto P(Q|D)P(D)$$

- The priori of document *D* is assumed to be uniform, which can be removed.
- $P(D)$ can also be assigned to some value that irrelevant to query, e.g., PageRank.



QLM

- QLM:

$$\begin{aligned} RSV(Q, D) &= P(Q \mid D) = P(Q \mid M_D) \\ &= P(q_1 q_2 \dots q_m \mid M_D) \\ &= P(q_1 \mid M_D) P(q_2 \mid M_D) \dots P(q_m \mid M_D) \\ &= \prod_{w \in Q} P(w \mid M_D)^{c(w, Q)} \end{aligned}$$

- The problem is transformed to the estimation of the unigram model M_D of document D, i.e., the term probability $P(w \mid M_D)$




Understanding QLM

- In QLM, $P(Q|D)$ is actually $P(Q|M_D)$. Thus, $P(Q|D)$ can not be referred to as the probability of generating Q from D .
- D and Q are both samples from the population distribution
- $P(w|M_D)$ or $P(w|D)$ can not be referred as the probability of w in D .



The process of QLM

- Step 1: based on D (sample), estimate document language model M_D (population). For the unigram model, i.e., calculate the term distribution $P(w|M_D)$.
- Step 2: calculate the likelihood of generating Q from M_D (generative probability)
- Step 3: rank documents based on the likelihood score



The estimation of M_D

- Problem: With known D (*sample*), estimate the parameters $P(w|M_D)$ in model M_D
- Maximum Likelihood Estimation, MLE
- MLE: maximize the likelihood of observed samples



MLE of M_D

- Assume the vocabulary size is L , then the parameters of M_D can be denoted as:

$$\begin{aligned}\vec{\theta}_D &= (\theta_1, \theta_2, \dots, \theta_L) \\ &= (P(w_1 | M_D), P(w_2 | M_D), \dots, P(w_L | M_D))\end{aligned}$$

- MLE:

$$\vec{\theta}_D^* = \arg \max_{\vec{\theta}_D} P(D | \vec{\theta}_D)$$

- The key is to calculate

$$P(D | \vec{\theta}_D)$$



An example of MLE estimation

- $D =$
($\langle I, 1 \rangle, \langle \text{like}, 1 \rangle, \langle \text{based}, 1 \rangle, \langle \text{statistical}, 1 \rangle, \langle \text{language}, 1 \rangle,$
 $\langle \text{model}, 2 \rangle, \langle \text{on}, 1 \rangle, \langle \text{information}, 1 \rangle, \langle \text{retrieval}, 1 \rangle$)

- Based on MLE estimation:

$$\begin{aligned} P(I|M_D) &= P(\text{like}|M_D) = P(\text{based}|M_D) \\ &= P(\text{statistical}|M_D) = P(\text{language}|M_D) = P(\text{on}|M_D) \\ &= P(\text{information}|M_D) = P(\text{retrieval}|M_D) = 0.1 \end{aligned}$$

$$P(\text{model}|M_D) = 0.2$$

all other probabilities are 0

The zero-probability problem in MLE



- For any term absent in D , its MLE probability is 0 (data sparseness). However, this term may present in new documents.
- Analogy: words that author does not use in one article may be used in future articles.

Several Smoothing techniques for QLM



- Jelinek-Mercer(JM), $0 \leq \lambda \leq 1$, mixture of document model and collection model

$$p(w|D) = \lambda p_{ML}(w|D) + (1 - \lambda) p(w|C)$$

- Dirichlet Priors(Dir), $\mu \geq 0$, *DIR entails JM and vice versa.*


$$p(w|D) = \frac{c(w,D) + \mu p(w|C)}{|D| + \mu}$$

- Absolute Discounting(Abs), $0 \leq \delta \leq 1$, $|D|_u$ denote the number of unique terms in D (u =unique)

$$p(w|D) = \frac{\max(c(w,D) - \delta, 0)}{|D|} + \frac{\delta |D|_u}{|D|} p(w|C)$$



Extensions for SLMIR



Three ways of using SLM in IR: a) query likelihood; b) document likelihood; c) model comparison



- **Neural Language Models for IR**

Problem with semantic similarity between words



枯燥

$$\begin{bmatrix} 0 \\ \mathbf{1} \\ 0 \\ \vdots \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

无聊

$$\begin{bmatrix} 0 \\ 0 \\ 0 \\ \vdots \\ 0 \\ \mathbf{1} \\ 0 \end{bmatrix}$$

枯燥 $\otimes$ 无聊

$$\begin{bmatrix} 0 \\ \mathbf{1} \\ 0 \\ \vdots \\ 0 \\ 0 \\ 0 \end{bmatrix}^T \times \begin{bmatrix} 0 \\ 0 \\ 0 \\ \vdots \\ 0 \\ \mathbf{1} \\ 0 \end{bmatrix} = 0$$




The similarity
between any different
words is 0




- In cognitive world, concepts are inter-connected and the semantics are distributed continuously



Continuous distributed representation of semantics



N-grams based on distributed representation



很

$$\begin{bmatrix} 0.01 \\ 0.59 \\ 0.18 \\ 0.05 \\ 0.47 \end{bmatrix}$$

$P(\text{无聊} |$

$$P \left(\begin{bmatrix} 0.25 \\ 0.12 \\ 0.39 \\ 0.46 \\ 0.26 \end{bmatrix} \middle| \begin{bmatrix} 0.01 \\ 0.59 \\ 0.18 \\ 0.05 \\ 0.47 \end{bmatrix} \right)$$

vs.

$P(\text{枯燥} |$


$$P \left(\begin{bmatrix} 0.24 \\ 0.15 \\ 0.42 \\ 0.51 \\ 0.21 \end{bmatrix} \middle| \begin{bmatrix} 0.01 \\ 0.59 \\ 0.18 \\ 0.05 \\ 0.47 \end{bmatrix} \right)$$

$$f \left(\begin{bmatrix} 0.01 \\ 0.59 \\ 0.18 \\ 0.05 \\ 0.47 \\ 0.25 \\ 0.12 \\ 0.39 \\ 0.46 \\ 0.26 \end{bmatrix} \right)$$

vs.

$$f \left(\begin{bmatrix} 0.01 \\ 0.59 \\ 0.18 \\ 0.05 \\ 0.47 \\ 0.24 \\ 0.15 \\ 0.42 \\ 0.51 \\ 0.21 \end{bmatrix} \right)$$

Deep Neural Network



Neural network language model




- Bengio2003 proposed Neural Network Language Model (NNLM)
- Multi-class prediction: given historical context, to predict which word should appear at the current position t
- Maximization:

$$L = \frac{1}{T} \sum_t \log f(w_t, w_{t-1}, \dots, w_{t-n+2}, w_{t-n+1}; \theta) + R(\theta)$$

- Word embedding representation as a by-product

$$P(y_t = i | \text{Context}) = P(y_t = i | w_{t-1}, \dots, w_{t-n+1})$$



Word2Vec



Continuous Bag-of-Words (CBOW)


- Given a context of a target word w_t (i.e., the c words before and c words after w_t), to predict what w_t is.

$$P(w_t | w_{t-c} : w_{t+c}) = \frac{\exp(\bar{v}^T v_{w_t})}{\sum_{n=1}^N \exp(\bar{v}^T v_n)}$$

$$\bar{v} = \frac{1}{2c} \sum_{-c \leq j \leq c, j \neq 0} v_j$$

- Maximization:**

$$L = \frac{1}{T} \sum_{t=1}^T \log P(w_t | w_{t-c} : w_{t+c})$$



Word2Vec




Skip-gram

- Given a target word w_t , to predict what its context are.

$$P(w_{t+j}|w_t) = \frac{\exp(w_{t+j}^T w_t)}{\sum_{n=1}^N \exp(w_{t+j}^T w_n)}$$

- Maximization:**



$$L = \frac{1}{T} \sum_{t=1}^T \sum_{-c \leq j \leq c, j \neq 0} \log P(w_{t+j}|w_t)$$




Word embedding vectors




- Property: $\text{vec}(\text{woman}) - \text{vec}(\text{man}) = \text{vec}(\text{queen}) - \text{vec}(\text{king})$



Neural IR models



MatchPyramid





Relevance Feedback

- **Relevance feedback:** to take the results that are initially returned from a given query, determine whether or not these results are truly relevant, and then use the information to perform a new query.
- **Query expansion:** to expand the search query with synonymous or relevant terms (from relevance feedback documents or other external resources) to improve retrieval performance in information retrieval.



Basic idea of relevance feedback

- ① User initiates a short query
- ② Search engine returns a list documents
- ③ User makes judgments the returned documents: either relevant or non-relevant
- ④ Search engine generate a new query to represent information need based on the user judgments. Hopefully the new query would perform better than the original one.
- ⑤ Search engine process the new query and return new results
- ⑥ The new results (ideally speaking) could have better precision and recall.
- ⑦ Iterate the feedback several times

Different kinds of Relevance Feedback



- **User Feedback or Explicit Feedback:** obtain the user's explicit relevance judgment indicating the relevance of a document.
- **Implicit Feedback:** inferred from users' interaction behavior, such as documents they do and do not select for viewing, the duration of time spent viewing a document, or page browsing or scrolling actions
- **Pseudo Feedback or Blind Feedback:** assume that the top "k" ranked documents are relevant. The user is not involved in any extended interaction.



Implicit relevant feedback



- Infer relevant judgment from **user behavior**
- The feedback could be inaccurate; the user is not necessarily involved to give explicit feedbacks.

Different kinds of implicit feedback



- **Mouse/keyboard movement:**
 - Click URL, add to favorite, copy/paste, page browsing, page skipping.
- **Eye movement**
 - Eye tracking can track user's eye movement
 - Fixation,

Eye tracking





Summary on implicit feedback

- **Pro:**
 - No need for user participation, easy to apply
 - User behavior could reflect user interests
- **Con:**
 - Require complex behavior analysis
 - The accuracy is not guaranteed
 - Sometimes need extra facilities.



Pseudo-relevance feedback (PRF)

- It automates the manual part of relevance feedback
- PRF algorithm
 1. do normal retrieval to find an initial set of most relevant documents
 2. assume that the top "k" ranked documents are relevant
 3. do relevance feedback as before under this assumption (e.g., Rocchio)
- **Works well on mean performance**
- **Due to the lack of user judgments, the accuracy is not guaranteed for all queries. Sometimes even hurt the performance badly.**
- **Lead to query drift problem when iterated several times.**



- **Evaluation of IR**



Why do we evaluate IR ?

- Different methods can be compared and analyzed effectively, and we could understand how different factors affect the IR system, so as to promote the continuous improvement of the research in this field.
- The goal of IR system is to return accurate results as fast and comprehensive as possible with minimum cost.



Evaluation of IR

- **Efficiency**
 - Time complexity
 - Space complexity
 - Response time
- **Effectiveness**
 - Precision
 - Recall
 - Ranking
- **Other**
 - Topic Coverage
 - Access time
 - Timeliness



How to evaluate?

- Same document collection, same query topics, same measurements, compare different IR systems.
 - **The Cranfield Experiments**, Cyril W. Cleverdon, 1957 –1968 (hundreds of documents)
 - **SMART System**, Gerald Salton, 1964-1988 (thousands of documents)
 - **TREC (Text REtrieval Conference)**, Donna Harman, National Institute of Standards and Technology, 1992 - (millions of documents), 'Olympics' in IR



Example for evaluation

Two IR systems, same set of queries, same document collection.


Which system is better?

Ranking	1	2	3	4	...
system1, query1	d3 ✓	d6 ✓	d8	d10	
system1, query2	d1	d4	d7	d11	
system2, query1	d6 ✓	d7	d2	d9 ✓	
system2, query2	d1	d2	d4	d13	

The ground truth for query 1 is {d3,d4,d6,d9}



Unranked retrieval evaluation





Unranked retrieval evaluation

- **Recall:** $RR/(RR + NR)$, fraction of relevant docs that are retrieved, $R \in [0,1]$
- **Precision:** $RR/(RR + RN)$, fraction of retrieved docs that are relevant, $P \in [0,1]$
- **Precision-Recall tradeoff:**
 - You can get high recall (but low precision) by retrieving all docs for all queries!
 - You can get high precision but low recall by only retrieving highly reliable results!



Back to example

system&query	1	2	3	4	5
system1, query1	d3 ✓	d6 ✓	d8	d10	d11
system1, query2	d1	d4	d7	d11	d13
system2, query1	d6 ✓	d7	d2	d9 ✓	/
system2, query2	d1	d2	d4	d13	d14

The ground truth for query 1 is {**d3,d4,d6,d9**}

For system 1, query 1, $P=2/5$, $R=2/4$

For system 2, query 1, $P=2/4$, $R=2/4$

Discussions in using precision/recall



- Both Precision and Recall are very important. Different applications and different users may have different demands for the two metrics.
 - spam filtering: can tolerate some spams, but require a low probability to misclassify normal emails.
 - Some prefers results with good coverage, since he has time to pick up useful information manually.
 - Someone may prefer results with high precision, since he does not need all information to complete the task.




A combined measure: F

- Combined measure that assesses precision/recall tradeoff is **F measure** (weighted harmonic mean):

$$F = \frac{1}{\alpha \frac{1}{P} + (1 - \alpha) \frac{1}{R}} = \frac{(\beta^2 + 1)PR}{\beta^2 P + R}$$

- People usually use balanced F_1 measure
 - i.e., with $\beta = 1$ or $\alpha = \frac{1}{2}$
- Harmonic mean is a conservative average
 - See CJ van Rijsbergen, *Information Retrieval*

Discussions in using precision/recall



- **Difficult to compare two systems since two metrics are only one-sided.**
 - Solution: combine two metrics into a single metric
- **Both metrics are based on unordered sets, thus ignore the ranking of the returned documents.**
 - Solution: take the ranking of the returned documents into account.
 - System that returns the relevant documents higher in the ranked lists is preferred.



Evaluating ranked results (1)

- **R-Precision**: The precision value obtained for the top R documents, where R is the ranking of the last relevant documents in the retrieved document list.

System1, query1	d3 ✓	d6 ✓	d8	d10	d11
System2, query1	d6 ✓	d7	d2	d9 ✓	

The ground truth for query 1 is {d3,d4,d6,d9}

$R-P1=2/2$ $R-P2=2/4$



Evaluating ranked results (2)

- precision versus recall curve :
 - The system can return any number of results
 - By taking various numbers of the top returned documents (at levels of recall), the evaluator can produce a *precision-recall curve*
 - *If the P-R curve of system 1 is on the above of that of system 2, then system 1 is better than system 2..*




Example of P-R curve

- The ground truth of query q:
 $R_q = \{d3, d5, d9, d25, d39, d44, d56, d71, d89, d123\}$
- The result of an IR system for query q is:


1. d123 R=0.1,P=1	6. d9 R=0.3,P=0.5	11. d38
2. d84	7. d511	12. d48
3. d56 R=0.2,P=0.67	8. d129	13. d250
4. d6	9. d187	14. d113
5. d8	10. d25 R=0.4,P=0.4	15. d3 R=0.5,P=0.33

P-R curve

Precision-recall 曲线



The Comparison of P-R curves for different systems






Interpolated precision

- In the previous example, assume $R_q = \{d3, d56, d129\}$
 - 3. d56 $R=0.33, P=0.33$; 8. d129 $R=0.66, P=0.25$;
15. d3 $R=1, P=0.2$
- There does not exist recall point at 10%, 20%, ..., 90%, but at 33.3%, 66.7%, 100%.
- Idea: take the max of precisions to right of value (interpolate)
- For recall = $t\%$, if there is no recall point, we define the precision to be the max value from recall $t\%$ to $(t+10)\%$.
- For example, the precisions at 0%, 10%, 20%, 30% are 0.33, and 40%~60% be 0.25, 70% be 0.2

P-R curve with Interpolated precision



Precision-Recall曲线





Evaluating ranked results(3)

- Average Precision, AP: the average of the precision at different recall point.
 - AP without interpolating: Assume query Q has 6 true relevant documents. IR system returned 5 of them, at position 1, 2, 5, 10, 20. Then, $AP = (1/1 + 2/2 + 3/5 + 4/10 + 5/20 + 0)/6$
 - AP with interpolating: average of the precision at recall point 0, 0.1, 0.2, ..., 1.0, equivalent to 11 point precisions
 - **Precision@N: the precision at position N.**

Evaluation metric over multiple queries



The way to calculate average:

- **Macro Average:** calculate certain metric for each query, and then calculate the arithmetic mean
- **MAP (Mean AP):** The macro average of AP for all queries
- **AvgP@N:** The macro average of Precision@N for all



Back to example

system&query	1	2	3	4	5
system1, query1	d3 ✓	d6 ✓	d8	d10	d11
system1, query2	d1 ✓	d4	d7	d11	d13 ✓
system2, query1	d6 ✓	d7	d2	d9 ✓	/
system2, query2	d1 ✓	d2 ✓	d4	d13 ✓	d14

The ground truth for query 1 and query 2 is {d3,d4,d6,d9}{d1,d2,d13}

system1, query1: $P@2=1$, $P@5=2/5$, $AP=(1+2/2+0+0)/4 = 1/2$;


system1, query2: $P@2=1/2$, $P@5=2/5$, $AP=(1/1+2/5+0)/3 = 7/15$;

System 1 MAP= $(1/2+7/15)/2 = 29/60 = 0.483$; AvgP@2=3/4; AvgP@5=2/5

system2, query1: $P@2=1/2$, $P@5=2/5$, $AP=(1+2/4+0+0)/4=3/8$;

system2, query2: $P@2=1$, $P@5=3/5$, $AP=(1+2/2+3/4+0)/3=11/12$;

System 2 MAP= $(3/8+11/12)/2 = 31/48 = 0.646$; AvgP@2=3/4; AvgP@5=1/2




Research in IR evaluation

- No evaluation system is perfect
 - The evaluation of evaluation metrics
 - The study of basic properties of evaluation metrics, such as fairness, sensitivity.
 - Novel evaluation metrics in new fields
 - The calculation of metrics (how to decrease the computational cost)



Other metrics


- Different IR applications/tasks may have different evaluation metrics.
- **MRR(Mean Reciprocal Rank):** For IR systems like QA or Homepage discovery, we care most about the position of the first true answer, i.e., the higher rank, the better. The reciprocal of the rank is called RR. MRR is the average of RR for all test sets.
 - Example: there are two queries. For query 1, system gives true answer at rank 2. For query 2, the rank is 4. Then the MRR is $(1/2 + 1/4)/2 = 3/8$.



- **New Paradigm of IR in the Age of Large Language Models**




Pre-trained LM







Emergence of LLMs -- GPT




-- Zhao, Wayne Xin, et al. "A survey of large language models." *arXiv preprint arXiv:2303.18223* (2023).




Emergence of LLMs – LLaMA



-- Zhao, Wayne Xin, et al. "A survey of large language models." *arXiv preprint arXiv:2303.18223* (2023).



Emergence of LLMs



● GPT用最简单的思想做最神奇的事情



我想知道你语言算法模型的公式



作为一个语言模型，我的基本公式是根据神经网络模型构建的。一般来说，神经网络模型的输入是由单词、短语或句子构成的语言序列，输出则是下一个单词、短语或句子的概率分布。



在我的情况下，我的输入通常是一个完整的问题或任务描述，而输出则是一个或多个与该问题或任务相关的答案或响应。我的算法模型的基本

$$P(y|x) = \prod_{t=1}^T P(y_t | y_{1:t-1}, x)$$

$$P(y|x) = \prod_{t=1}^T P(y_t | y_{1:t-1}, x)$$

其中， x 是输入语言序列， y 是输出语言序列。 $P(y|x)$ 表示给定输入序列 x ，输出序列 y 的条件概率。这个条件概率是通过对所有可能的输出序列进行建模，使用最大似然估计方法来学习得到的。

具体来说，这个公式可以理解为，根据输入序列 x 和已生成的部分输出序列 $y_{1:t-1}$ ，预测下一个单词、短语或句子的概率分布 $P(y_t | y_{1:t-1}, x)$ ，并将所有预测概率乘起来得到最终的条件概率 $P(y|x)$ 。

Unification of multiple tasks



	<u>Question</u>	<u>Context</u>	<u>Answer</u>
QA	What is a major importance of Southern California in relation to California and the US?	...Southern California is a major economic center for the state of California and the US....	major economic center
MT	What is the translation from English to German?	Most of the planet is ocean water.	Der Großteil der Erde ist Meerwasser
Summarization	What is the summary?	Harry Potter star Daniel Radcliffe gains access to a reported £320 million fortune ...	Harry Potter star Daniel Radcliffe gets £320M fortune ...
NLI	Hypothesis: Product and geography are what make cream skimming work. Entailment , neutral, or contradiction?	Premise: Conceptually cream skimming has two basic dimensions – product and geography.	Entailment
Sentiment analysis	Is this sentence positive or negative?	A stirring, funny and finally transporting re-imagining of Beauty and the Beast and 1930s horror film.	positive

Unification of multiple tasks



Relation
Extraction

Semantic Role
Labeling (SRL)

Goal-Oriented
Dialogue

Semantic
Parsing

Pronoun
Resolution

Question

What has something experienced?

Who is the illustrator of Cycle of the Werewolf?

What is the change in dialogue state?

What is the translation from English to SQL?

Who had given help?
Susan or Joan?

Context

Areas of the Baltic that have experienced eutrophication.

Cycle of the Werewolf is a short novel by Stephen King, featuring illustrations by comic book artist Bernie Wrightson.

Are there any Eritrean restaurants in town?

The table has column names...
Tell me what the notes are for South Australia

Joan made sure to thank Susan for all the help she had given.

Answer

eutrophication

Bernie Wrightson

food: Eritrean

SELECT notes from table
WHERE
'Current Slogan' =
'South Australia'

Susan

Prompt to LLM



Zero-shot prompt

Here is a goal: Get a cool sleep on summer days.

How would you accomplish this goal?

OPTIONS:

-Keep stack of pillow cases in fridge.

-Keep stack of pillow cases in oven.

keep stack of pillow cases in fridge

Give some examples (optinal)

Argument 1: " We' ve been spending a lot of time in Los Angeles talking to TV production people "


Argument 2: "With the competitiveness of the television market these days, everyone is looking for a way to get viewers more excited "

Question:What is the discourse relation between Argument 1 and Argument 2?

- A. Comparison
- B. Contingency
- C. Expansion
- D. Temporal

Let's think step by step(optional)

Prompt to LLM



Few-shot prompt

Model Input

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: The answer is 11.

Q: The cafeteria had 23 apples. If they used 20 to make lunch and bought 6 more, how many apples do they have?

Model Output

A: The answer is 27.

Few shot Chain of thought prompt

Model Input

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: Roger started with 5 balls. 2 cans of 3 tennis balls each is 6 tennis balls. $5 + 6 = 11$. The answer is 11.

Q: The cafeteria had 23 apples. If they used 20 to make lunch and bought 6 more, how many apples do they have?

Model Output

A: The cafeteria had 23 apples originally. They used 20 to make lunch. So they had $23 - 20 = 3$. They bought 6 more apples, so they have $3 + 6 = 9$. The answer is 9.

Prompt to LLM



- Zero-shot CoT Prompting

(a) Few-shot

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: The answer is 11.

Q: A juggler can juggle 16 balls. Half of the balls are golf balls, and half of the golf balls are blue. How many blue golf balls are there?

A:

(Output) The answer is 8. ✗

(b) Few-shot-CoT

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: Roger started with 5 balls. 2 cans of 3 tennis balls each is 6 tennis balls. $5 + 6 = 11$. The answer is 11.

Q: A juggler can juggle 16 balls. Half of the balls are golf balls, and half of the golf balls are blue. How many blue golf balls are there?

A:

(Output) The juggler can juggle 16 balls. Half of the balls are golf balls. So there are $16 / 2 = 8$ golf balls. Half of the golf balls are blue. So there are $8 / 2 = 4$ blue golf balls. The answer is 4. ✓

(c) Zero-shot

Q: A juggler can juggle 16 balls. Half of the balls are golf balls, and half of the golf balls are blue. How many blue golf balls are there?

A: The answer (arabic numerals) is

(Output) 8 ✗

(d) Zero-shot-CoT (Ours)

Q: A juggler can juggle 16 balls. Half of the balls are golf balls, and half of the golf balls are blue. How many blue golf balls are there?

A: **Let's think step by step.**

(Output) There are 16 balls in total. Half of the balls are golf balls. That means that there are 8 golf balls. Half of the golf balls are blue. That means that there are 4 blue golf balls. ✓

--Kojima, Takeshi, et al. "Large language models are zero-shot reasoners." *Advances in neural information processing systems* 35 (2022): 22199-22213.



Summary

- **The importance of IR**
- **Basic IR concepts and models**
- **IR Evaluation**